

```
In [1]: import warnings
warnings.filterwarnings("ignore")
from sklearn.datasets import load_boston
from random import seed
from random import randrange
from csv import reader
from math import sqrt
from sklearn import preprocessing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from prettytable import PrettyTable
from sklearn.linear_model import SGDRegressor
from sklearn import preprocessing
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
```

```
In [2]: # Loading data into X,Y
X = load_boston().data
Y = load_boston().target
```

```
In [3]: #splitting data into X_train and X_test
import sklearn
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.33, random_state = 5)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

```
(339, 13)
(167, 13)
(339,)
(167,)
```

```
In [8]: #standardizing the data
scaler = preprocessing.StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```


```
In [9]: #clf = SGDRegressor(loss='squared_loss')
#clf.fit(X, Y)
#print(mean_squared_error(Y, clf.predict(X)))
```

```
In [12]: #Applying SGD regressor on Train data
clf = SGDRegressor(loss='squared_loss')
clf.fit(X_train, Y_train)
```

```
Out[12]: SGDRegressor(alpha=0.0001, average=False, epsilon=0.1, eta0=0.01,
    fit_intercept=True, l1_ratio=0.15, learning_rate='invscaling',
    loss='squared_loss', max_iter=None, n_iter=None, penalty='l2',
    power_t=0.25, random_state=None, shuffle=True, tol=None, verbose=0,
    warm_start=False)
```

```
In [13]: clf.fit(X_train, Y_train)
#Finding the predicted prices
Y_pred_clf = clf.predict(X_test)

#plotting a chart of actual prices vs predicted prices for SGD regressor
plt.scatter(Y_test, Y_pred_clf)
plt.xlabel("Prices:  $Y_i$ ")
plt.ylabel("Predicted prices:  $\hat{Y}_i$ ")
plt.title("Prices vs Predicted prices:  $Y_i$  vs  $\hat{Y}_i$ ")
plt.show()
```




```
In [14]: from sklearn.linear_model import LinearRegression

lm = LinearRegression()
lm.fit(X_train, Y_train)

#Finding the predicted price
Y_pred_lm = lm.predict(X_test)

#plotting a chart of actual prices vs predicted prices for linear regression
plt.scatter(Y_test, Y_pred_lm)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")
plt.show()
```



```
In [15]: #Finding weight vector of SGD
clf.coef_
```

```
Out[15]: array([-0.81363518,  0.40405972, -0.51748626,  0.29339603, -0.22368237,
                3.14507661, -0.42221585, -1.68006996,  0.74576312, -0.52451597,
                -1.77986047,  0.87611157, -2.98815837])
```

```
In [16]: #Finding weight vector of Linear regression
lm.coef_
```

```
Out[16]: array([-1.31193031,  0.86187745, -0.16719287,  0.18957843, -1.48658584,
                2.79131565, -0.32737703, -2.77204093,  2.97567549, -2.2727549 ,
                -2.13375869,  1.05842993, -3.33495407])
```

In [17]: [#https://stackoverflow.com/questions/41140647/python-printing-lists-with-tabulate](https://stackoverflow.com/questions/41140647/python-printing-lists-with-tabulate)

```
#Tabulating and comparing SGD and Linear regression weights
from tabulate import tabulate
headers = ['SGD Weights','Linear regression weights']

table = zip(clf.coef_,lm.coef_)
print(tabulate(table, headers=headers, floatfmt=".4f"))
```

| SGD Weights | Linear regression weights |
|-------------|---------------------------|
| -----       | -----                     |
| -0.8136     | -1.3119                   |
| 0.4041      | 0.8619                    |
| -0.5175     | -0.1672                   |
| 0.2934      | 0.1896                    |
| -0.2237     | -1.4866                   |
| 3.1451      | 2.7913                    |
| -0.4222     | -0.3274                   |
| -1.6801     | -2.7720                   |
| 0.7458      | 2.9757                    |
| -0.5245     | -2.2728                   |
| -1.7799     | -2.1338                   |
| 0.8761      | 1.0584                    |
| -2.9882     | -3.3350                   |

In [18]: `clf.fit(X_train, Y_train)`  
`print(mean_squared_error(Y_train, clf.predict(X_train)))`  
21.39072011815044

In [19]: `lm.fit(X_train, Y_train)`  
`print(mean_squared_error(Y_train, lm.predict(X_train)))`  
19.54675847353467

In [20]: *#calculating mean squared errors of both models*  
`a=(Y_test-clf.predict(X_test))`  
`b=(Y_test-lm.predict(X_test))`  
  
`a=a**2`  
`b=b**2`

```
In [21]: #comparing MSE of SGD regressor and Linear regression  
from tabulate import tabulate  
headers = ['SGD MSE','Linear regression MSE']  
  
table = zip(a,b)  
print(tabulate(table, headers=headers, floatfmt=".4f"))
```

| SGD MSE  | Linear regression MSE |
|----------|-----------------------|
| 0.6603   | 0.0176                |
| 1.0227   | 12.1909               |
| 16.9378  | 20.4322               |
| 100.6309 | 53.7519               |
| 3.8395   | 2.4659                |
| 38.3353  | 22.3665               |
| 8.0225   | 9.8563                |
| 0.0056   | 0.8594                |
| 13.7791  | 12.6495               |
| 8.6762   | 0.5588                |
| 0.6498   | 0.9899                |
| 4.7137   | 7.7641                |
| 5.0103   | 9.6807                |
| 19.3146  | 6.2054                |
| 5.5706   | 15.1025               |
| 0.6116   | 4.2494                |
| 410.3948 | 309.7293              |
| 0.0879   | 0.7770                |
| 30.9380  | 27.3459               |
| 28.5613  | 15.1458               |
| 13.6688  | 29.2147               |
| 13.4311  | 14.3116               |
| 50.4955  | 40.8117               |
| 0.9498   | 1.6113                |
| 0.0909   | 0.3423                |
| 0.2017   | 0.3870                |
| 2.1558   | 3.3439                |
| 3.9715   | 4.4596                |
| 1.3825   | 0.3322                |
| 0.1089   | 11.8969               |
| 17.7918  | 1.5997                |
| 3.1109   | 0.0064                |
| 9.4612   | 24.5529               |
| 21.3465  | 25.1473               |
| 3.1443   | 7.9492                |
| 0.7905   | 0.9731                |
| 6.6289   | 10.7388               |
| 0.5360   | 0.2552                |
| 48.8440  | 51.1024               |
| 5.0086   | 0.7585                |
| 20.2667  | 40.7924               |
| 123.2739 | 109.0770              |
| 3.4386   | 0.1671                |
| 2.4091   | 0.2696                |
| 0.0134   | 0.0088                |
| 42.5192  | 34.0705               |
| 195.8100 | 175.5499              |
| 12.8017  | 0.3066                |
| 1.3089   | 0.0026                |
| 0.5680   | 0.4035                |
| 0.5155   | 0.0001                |
| 4.4839   | 5.4124                |
| 3.8192   | 4.7375                |
| 0.0911   | 0.0097                |
| 0.5679   | 3.0117                |

|          |          |
|----------|----------|
| 2.5984   | 0.0055   |
| 65.8017  | 71.9992  |
| 0.3204   | 1.9132   |
| 0.0673   | 0.0085   |
| 3.8915   | 4.1037   |
| 12.0519  | 14.2992  |
| 4.4805   | 8.3210   |
| 4.9402   | 2.8286   |
| 94.9936  | 79.6567  |
| 7.1728   | 3.5752   |
| 4.5771   | 15.1269  |
| 3.2892   | 1.1141   |
| 0.1442   | 0.2851   |
| 0.2992   | 0.0053   |
| 2.9018   | 0.4932   |
| 3.2614   | 15.0716  |
| 11.8745  | 21.5886  |
| 3.5886   | 3.4967   |
| 115.1919 | 77.1025  |
| 136.8576 | 78.6721  |
| 8.7710   | 18.7876  |
| 13.3833  | 20.0324  |
| 1.6009   | 0.0954   |
| 12.8608  | 17.0084  |
| 4.5382   | 14.7832  |
| 0.0359   | 8.0436   |
| 104.5661 | 112.3588 |
| 0.6322   | 0.3682   |
| 5.4019   | 10.7180  |
| 23.8102  | 34.7275  |
| 10.2004  | 18.6629  |
| 0.1739   | 1.1758   |
| 3.4962   | 0.1004   |
| 0.8049   | 0.4569   |
| 3.9801   | 1.8368   |
| 203.0105 | 188.7152 |
| 3.6719   | 4.1783   |
| 6.6290   | 0.0850   |
| 0.4778   | 0.4683   |
| 15.3659  | 4.5165   |
| 9.8774   | 9.5345   |
| 11.0441  | 11.1907  |
| 0.5707   | 17.5903  |
| 6.6314   | 1.9465   |
| 1.7963   | 0.1687   |
| 5.4615   | 0.0084   |
| 4.1824   | 9.0657   |
| 1.0552   | 1.1016   |
| 67.8689  | 52.4643  |
| 35.3341  | 46.0920  |
| 0.4646   | 1.9360   |
| 21.1132  | 38.7513  |
| 989.9258 | 751.1242 |
| 20.6970  | 23.3340  |
| 0.4032   | 5.7133   |
| 16.8595  | 13.8040  |
| 1.3212   | 0.1264   |

|          |          |
|----------|----------|
| 30.8442  | 52.2018  |
| 10.3347  | 5.6124   |
| 12.3023  | 20.5187  |
| 0.3668   | 0.9498   |
| 30.9493  | 13.7672  |
| 2.9821   | 2.6733   |
| 2.5935   | 0.2543   |
| 0.6682   | 1.9840   |
| 5.9755   | 9.4744   |
| 0.0890   | 4.4263   |
| 4.7010   | 2.3728   |
| 14.2292  | 9.6116   |
| 0.8943   | 1.3058   |
| 1.3961   | 1.4449   |
| 5.8583   | 0.8060   |
| 0.0174   | 2.2955   |
| 15.3988  | 17.5098  |
| 100.3884 | 82.4811  |
| 150.7083 | 147.6452 |
| 10.7323  | 20.6138  |
| 1.2001   | 4.1400   |
| 4.6595   | 3.5877   |
| 33.5495  | 54.0743  |
| 0.2894   | 0.3950   |
| 0.7210   | 3.9191   |
| 48.2809  | 23.7523  |
| 3.0517   | 3.2417   |
| 0.5713   | 0.0166   |
| 7.0136   | 13.0412  |
| 18.5664  | 7.2417   |
| 19.1313  | 16.0424  |
| 14.9465  | 5.3290   |
| 0.9669   | 0.0136   |
| 867.0612 | 707.2908 |
| 47.9280  | 37.6687  |
| 0.8946   | 0.3658   |
| 4.1980   | 5.6095   |
| 47.4987  | 43.4578  |
| 0.1041   | 0.0038   |
| 431.5750 | 329.0273 |
| 0.2202   | 1.8615   |
| 1.9557   | 8.5125   |
| 0.3472   | 7.9701   |
| 17.1837  | 16.0415  |
| 9.2588   | 20.8928  |
| 0.7249   | 0.0824   |
| 12.4961  | 19.7446  |
| 0.2006   | 0.5149   |
| 3.3543   | 6.1065   |
| 14.4457  | 24.5683  |
| 20.5040  | 21.9177  |
| 3.1268   | 0.4827   |
| 0.0327   | 0.1633   |
| 2.4052   | 3.4385   |
| 0.4784   | 1.4528   |



# Conclusions

1. Mean squared error of SGD regressor is higher than that of linear regression
2. Most of the SGD weights are higher than that of linear model's weights