

▼ 3-Convolution Layers

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(128, kernel_size=(2, 2), activation='relu', input_shape=input_shape, kernel_
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Conv2D(32, (5, 5), activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.Adam())

model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



Using TensorFlow backend.

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>

11493376/11490434 [=====] - 2s 0us/step

(60000, 28, 28)

x_train shape: (60000, 28, 28, 1)

60000 train samples

10000 test samples

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 27, 27, 128)	640
batch_normalization_1 (Batch Normalization)	(None, 27, 27, 128)	512
conv2d_2 (Conv2D)	(None, 25, 25, 64)	73792
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
conv2d_3 (Conv2D)	(None, 8, 8, 32)	51232
batch_normalization_2 (Batch Normalization)	(None, 8, 8, 32)	128
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 32)	0
dropout_1 (Dropout)	(None, 4, 4, 32)	0
flatten_1 (Flatten)	(None, 512)	0
dense_1 (Dense)	(None, 128)	65664
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290

=====
Total params: 193,258

Trainable params: 192,938

Non-trainable params: 320

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:79

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 22s 374us/step - loss: 0.2054 - acc: 0

Epoch 2/12

60000/60000 [=====] - 19s 310us/step - loss: 0.0615 - acc: 0

Epoch 3/12

60000/60000 [=====] - 19s 310us/step - loss: 0.0451 - acc: 0

Epoch 4/12

60000/60000 [=====] - 19s 311us/step - loss: 0.0395 - acc: 0

Epoch 5/12

60000/60000 [=====] - 19s 311us/step - loss: 0.0314 - acc: 0

Epoch 6/12

60000/60000 [=====] - 19s 312us/step - loss: 0.0287 - acc: 0

Epoch 7/12

60000/60000 [=====] - 19s 311us/step - loss: 0.0268 - acc: 0

Epoch 8/12

60000/60000 [=====] - 19s 310us/step - loss: 0.0243 - acc: 0

Epoch 9/12

60000/60000 [=====] - 19s 310us/step - loss: 0.0218 - acc: 0

Epoch 10/12

60000/60000 [=====] - 19s 310us/step - loss: 0.0207 - acc: 0

Epoch 11/12

60000/60000 [=====] - 19s 311us/step - loss: 0.0199 - acc: 0

Epoch 12/12

60000/60000 [=====] - 19s 311us/step - loss: 0.0176 - acc: 0

Test loss: 0.028177190259779036

Test accuracy: 0.9919

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
```

```
batch_size = 128
num_classes = 10
epochs = 12
```

```
# input image dimensions
img_rows, img_cols = 28, 28
```

```
# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
print(x_train.shape)
```

```
if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
```

```
else:
```

```
x_test.
```

```
x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(128, kernel_size=(2, 2), activation='relu', input_shape=input_shape, kernel_
model.add(BatchNormalization()))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3), activation='sigmoid'))
model.add(Conv2D(32, (5, 5), activation='sigmoid'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(128, activation='sigmoid'))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.Adagrad

model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 27, 27, 128)	640
batch_normalization_3 (Batch Normalization)	(None, 27, 27, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 13, 13, 128)	0
conv2d_5 (Conv2D)	(None, 11, 11, 64)	73792
conv2d_6 (Conv2D)	(None, 7, 7, 32)	51232
batch_normalization_4 (Batch Normalization)	(None, 7, 7, 32)	128
max_pooling2d_4 (MaxPooling2D)	(None, 3, 3, 32)	0
flatten_2 (Flatten)	(None, 288)	0
dense_3 (Dense)	(None, 128)	36992
dense_4 (Dense)	(None, 10)	1290

```
=====  
Total params: 164,586  
Trainable params: 164,266  
Non-trainable params: 320
```

```
=====  
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/12  
60000/60000 [=====] - 14s 233us/step - loss: 0.1210 - acc: 0  
Epoch 2/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0365 - acc: 0  
Epoch 3/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0247 - acc: 0  
Epoch 4/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0185 - acc: 0  
Epoch 5/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0141 - acc: 0  
Epoch 6/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0107 - acc: 0  
Epoch 7/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0086 - acc: 0  
Epoch 8/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0070 - acc: 0  
Epoch 9/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0058 - acc: 0  
Epoch 10/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0049 - acc: 0  
Epoch 11/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0041 - acc: 0  
Epoch 12/12  
60000/60000 [=====] - 13s 220us/step - loss: 0.0036 - acc: 0  
Test loss: 0.023953616692498327  
Test accuracy: 0.993
```

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(128, kernel_size=(2, 2), activation='tanh', input_shape=input_shape, kernel_
model.add(Conv2D(64, (3, 3), activation='tanh'))
model.add(Dropout(0.5))
model.add(Conv2D(32, (5, 5), activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))

```

```
model.add(Dense(100, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimizers.Adadelta)

model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, verbose=1, validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```




```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 27, 27, 128)	640
conv2d_8 (Conv2D)	(None, 25, 25, 64)	73792
dropout_3 (Dropout)	(None, 25, 25, 64)	0
conv2d_9 (Conv2D)	(None, 21, 21, 32)	51232
max_pooling2d_5 (MaxPooling2D)	(None, 10, 10, 32)	0
flatten_3 (Flatten)	(None, 3200)	0
dense_5 (Dense)	(None, 128)	409728
dropout_4 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 10)	1290

```
=====
Total params: 536,682
Trainable params: 536,682
Non-trainable params: 0
```

```
=====
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/12
60000/60000 [=====] - 19s 316us/step - loss: 0.2839 - acc: 0
Epoch 2/12
60000/60000 [=====] - 18s 303us/step - loss: 0.1130 - acc: 0
Epoch 3/12
60000/60000 [=====] - 18s 302us/step - loss: 0.0850 - acc: 0
Epoch 4/12
60000/60000 [=====] - 18s 303us/step - loss: 0.0677 - acc: 0
Epoch 5/12
60000/60000 [=====] - 18s 303us/step - loss: 0.0577 - acc: 0
Epoch 6/12
60000/60000 [=====] - 18s 303us/step - loss: 0.0522 - acc: 0
Epoch 7/12
60000/60000 [=====] - 18s 302us/step - loss: 0.0453 - acc: 0
Epoch 8/12
60000/60000 [=====] - 18s 303us/step - loss: 0.0417 - acc: 0
Epoch 9/12
60000/60000 [=====] - 18s 305us/step - loss: 0.0383 - acc: 0
Epoch 10/12
60000/60000 [=====] - 18s 303us/step - loss: 0.0331 - acc: 0
Epoch 11/12
60000/60000 [=====] - 18s 302us/step - loss: 0.0297 - acc: 0
Epoch 12/12
60000/60000 [=====] - 18s 302us/step - loss: 0.0289 - acc: 0
Test loss: 0.02994774283710726
Test accuracy: 0.9904
```

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(128, kernel_size=(2, 2), activation='relu', input_shape=input_shape, kernel_
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(Conv2D(32, (5, 5), activation='relu'))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))

model.summary()
```

```

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.sgd(),

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

```

↳ (60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_4"

```

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 27, 27, 128)	640
conv2d_11 (Conv2D)	(None, 25, 25, 64)	73792
conv2d_12 (Conv2D)	(None, 21, 21, 32)	51232
flatten_4 (Flatten)	(None, 14112)	0
dense_7 (Dense)	(None, 128)	1806464
dense_8 (Dense)	(None, 10)	1290

```

Total params: 1,933,418
Trainable params: 1,933,418
Non-trainable params: 0

```

```

Train on 60000 samples, validate on 10000 samples

```

```

Epoch 1/12
60000/60000 [=====] - 17s 276us/step - loss: 0.6677 - acc: 0
Epoch 2/12
60000/60000 [=====] - 16s 266us/step - loss: 0.2070 - acc: 0
Epoch 3/12
60000/60000 [=====] - 16s 267us/step - loss: 0.1461 - acc: 0
Epoch 4/12
60000/60000 [=====] - 16s 266us/step - loss: 0.1135 - acc: 0
Epoch 5/12
60000/60000 [=====] - 16s 266us/step - loss: 0.0935 - acc: 0
Epoch 6/12
60000/60000 [=====] - 16s 267us/step - loss: 0.0785 - acc: 0
Epoch 7/12
60000/60000 [=====] - 16s 267us/step - loss: 0.0676 - acc: 0
Epoch 8/12
60000/60000 [=====] - 16s 266us/step - loss: 0.0576 - acc: 0
Epoch 9/12
60000/60000 [=====] - 16s 266us/step - loss: 0.0502 - acc: 0
Epoch 10/12
60000/60000 [=====] - 16s 267us/step - loss: 0.0447 - acc: 0
Epoch 11/12
60000/60000 [=====] - 16s 266us/step - loss: 0.0402 - acc: 0
Epoch 12/12
60000/60000 [=====] - 16s 266us/step - loss: 0.0360 - acc: 0
Test loss: 0.06601626948665362
Test accuracy: 0.979

```

▼ 5-Layers

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 512
num_classes = 10
epochs = 10

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(128, kernel_size=(2, 2), activation='relu', input_shape=input_shape, kernel_
model.add(Conv2D(96, (3, 3), activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
```

```
model.add(Conv2D(84, (5, 5),padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Conv2D(32, (7, 7),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam())

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_5"
```

Layer (type)	Output Shape	Param #
conv2d_13 (Conv2D)	(None, 27, 27, 128)	640
conv2d_14 (Conv2D)	(None, 25, 25, 96)	110688
batch_normalization_5 (Batch Normalization)	(None, 25, 25, 96)	384
max_pooling2d_6 (MaxPooling2D)	(None, 12, 12, 96)	0
dropout_5 (Dropout)	(None, 12, 12, 96)	0
conv2d_15 (Conv2D)	(None, 12, 12, 84)	201684
batch_normalization_6 (Batch Normalization)	(None, 12, 12, 84)	336
conv2d_16 (Conv2D)	(None, 12, 12, 64)	48448
max_pooling2d_7 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_6 (Dropout)	(None, 6, 6, 64)	0
conv2d_17 (Conv2D)	(None, 6, 6, 32)	100384
max_pooling2d_8 (MaxPooling2D)	(None, 3, 3, 32)	0
dropout_7 (Dropout)	(None, 3, 3, 32)	0
flatten_5 (Flatten)	(None, 288)	0
dense_9 (Dense)	(None, 128)	36992
dropout_8 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 10)	1290

```
=====
Total params: 500,846
Trainable params: 500,486
Non-trainable params: 360
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/10
60000/60000 [=====] - 26s 430us/step - loss: 0.7600 - acc: 0
Epoch 2/10
60000/60000 [=====] - 23s 377us/step - loss: 0.1271 - acc: 0
Epoch 3/10
60000/60000 [=====] - 23s 376us/step - loss: 0.0804 - acc: 0
Epoch 4/10
60000/60000 [=====] - 23s 376us/step - loss: 0.0699 - acc: 0
Epoch 5/10
60000/60000 [=====] - 23s 376us/step - loss: 0.0552 - acc: 0
Epoch 6/10
60000/60000 [=====] - 23s 377us/step - loss: 0.0461 - acc: 0
Epoch 7/10
60000/60000 [=====] - 23s 375us/step - loss: 0.0419 - acc: 0
```

```

Epoch 8/10
60000/60000 [=====] - 23s 376us/step - loss: 0.0383 - acc: 0
Epoch 9/10
60000/60000 [=====] - 23s 376us/step - loss: 0.0346 - acc: 0
Epoch 10/10
60000/60000 [=====] - 23s 376us/step - loss: 0.0296 - acc: 0
Test loss: 0.02720410974002425
Test accuracy: 0.9921

```

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 512
num_classes = 10
epochs = 10

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(128, kernel_size=(2, 2), activation='relu', input_shape=input_shape, kernel

```

```
model.add(Conv2D(128, (3, 3), kernel_size=(2, 2), activation='relu', input_shape=input_shape, kernel_size=
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(96, (3, 3), activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())

model.add(Conv2D(84, (5, 5),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(32, (7, 7),padding='same', activation='relu'))

model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam())

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```




```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_6"
```

Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 27, 27, 128)	640
batch_normalization_7 (Batch Normalization)	(None, 27, 27, 128)	512
max_pooling2d_9 (MaxPooling2D)	(None, 13, 13, 128)	0
conv2d_19 (Conv2D)	(None, 11, 11, 96)	110688
batch_normalization_8 (Batch Normalization)	(None, 11, 11, 96)	384
max_pooling2d_10 (MaxPooling2D)	(None, 5, 5, 96)	0
batch_normalization_9 (Batch Normalization)	(None, 5, 5, 96)	384
conv2d_20 (Conv2D)	(None, 5, 5, 84)	201684
max_pooling2d_11 (MaxPooling2D)	(None, 2, 2, 84)	0
batch_normalization_10 (Batch Normalization)	(None, 2, 2, 84)	336
conv2d_21 (Conv2D)	(None, 2, 2, 64)	48448
max_pooling2d_12 (MaxPooling2D)	(None, 1, 1, 64)	0
conv2d_22 (Conv2D)	(None, 1, 1, 32)	100384
flatten_6 (Flatten)	(None, 32)	0
dense_11 (Dense)	(None, 128)	4224
dense_12 (Dense)	(None, 10)	1290
Total params: 468,974		
Trainable params: 468,166		
Non-trainable params: 808		

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 17s 278us/step - loss: 0.4008 - acc: 0

Epoch 2/10

60000/60000 [=====] - 15s 244us/step - loss: 0.0344 - acc: 0

Epoch 3/10

60000/60000 [=====] - 15s 244us/step - loss: 0.0218 - acc: 0

Epoch 4/10

60000/60000 [=====] - 15s 244us/step - loss: 0.0127 - acc: 0

Epoch 5/10

60000/60000 [=====] - 15s 244us/step - loss: 0.0077 - acc: 0

Epoch 6/10

60000/60000 [=====] - 15s 243us/step - loss: 0.0070 - acc: 0

Epoch 7/10

60000/60000 [=====] - 15s 243us/step - loss: 0.0045 - acc: 0

Epoch 8/10

60000/60000 [=====] - 15s 244us/step - loss: 0.0050 - acc: 0

```

Epoch 9/10
60000/60000 [=====] - 15s 245us/step - loss: 0.0061 - acc: 0
Epoch 10/10
60000/60000 [=====] - 15s 243us/step - loss: 0.0056 - acc: 0
Test loss: 0.04050913316059923
Test accuracy: 0.9901

```

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 512
num_classes = 10
epochs = 10

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(512, kernel_size=(2, 2), activation='sigmoid', input_shape=input_shape, kern
model.add(Conv2D(256, (3, 3), padding='same', activation='sigmoid'))
model.add(Conv2D(128, (5, 5), padding='same', activation='sigmoid'))

```

```
model.add(Conv2D(128, (5, 5),padding= 'same' , activation= sigmoid ))

model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.5))


model.add(Conv2D(64, (3, 3),padding='same', activation='sigmoid'))
model.add(Conv2D(32, (7, 7),padding='same', activation='sigmoid'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.5))


model.add(Dropout(0.5))
model.add(Flatten())
model.add(Dense(128, activation='sigmoid'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adagrad

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_7"
```

Layer (type)	Output Shape	Param #
conv2d_23 (Conv2D)	(None, 27, 27, 512)	2560
conv2d_24 (Conv2D)	(None, 27, 27, 256)	1179904
conv2d_25 (Conv2D)	(None, 27, 27, 128)	819328
max_pooling2d_13 (MaxPooling)	(None, 13, 13, 128)	0
dropout_9 (Dropout)	(None, 13, 13, 128)	0
conv2d_26 (Conv2D)	(None, 13, 13, 64)	73792
conv2d_27 (Conv2D)	(None, 13, 13, 32)	100384
max_pooling2d_14 (MaxPooling)	(None, 6, 6, 32)	0
dropout_10 (Dropout)	(None, 6, 6, 32)	0
dropout_11 (Dropout)	(None, 6, 6, 32)	0
flatten_7 (Flatten)	(None, 1152)	0
dense_13 (Dense)	(None, 128)	147584
dropout_12 (Dropout)	(None, 128)	0
dense_14 (Dense)	(None, 10)	1290
Total params: 2,324,842		
Trainable params: 2,324,842		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/10
60000/60000 [=====] - 108s 2ms/step - loss: 2.3278 - acc: 0.
Epoch 2/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3027 - acc: 0.1
Epoch 3/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3020 - acc: 0.1
Epoch 4/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3016 - acc: 0.1
Epoch 5/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3015 - acc: 0.1
Epoch 6/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3017 - acc: 0.1
Epoch 7/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3014 - acc: 0.1
Epoch 8/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3015 - acc: 0.1
Epoch 9/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3012 - acc: 0.1
Epoch 10/10
60000/60000 [=====] - 94s 2ms/step - loss: 2.3012 - acc: 0.1
```

Test loss: 2.301075766372681

Test accuracy: 0.1135

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 512
num_classes = 10
epochs = 10

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(256, kernel_size=(2, 2),padding='same',activation='tanh',input_shape=input_shape))
model.add(Conv2D(128, (3, 3),padding='same', activation='tanh'))
model.add(Conv2D(128, (3, 3),padding='same', activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Conv2D(64, (5, 5),padding='same', activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(32, (7, 7),padding='same', activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(128, activation='tanh'))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adadelta)

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_8"
```

Layer (type)	Output Shape	Param #
conv2d_28 (Conv2D)	(None, 28, 28, 256)	1280
conv2d_29 (Conv2D)	(None, 28, 28, 128)	295040
conv2d_30 (Conv2D)	(None, 28, 28, 128)	147584
max_pooling2d_15 (MaxPooling)	(None, 14, 14, 128)	0
conv2d_31 (Conv2D)	(None, 14, 14, 64)	204864
max_pooling2d_16 (MaxPooling)	(None, 7, 7, 64)	0
conv2d_32 (Conv2D)	(None, 7, 7, 32)	100384
max_pooling2d_17 (MaxPooling)	(None, 3, 3, 32)	0
flatten_8 (Flatten)	(None, 288)	0
dense_15 (Dense)	(None, 128)	36992
dense_16 (Dense)	(None, 10)	1290
Total params: 787,434		
Trainable params: 787,434		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/10
60000/60000 [=====] - 50s 830us/step - loss: 0.3525 - acc: 0
Epoch 2/10
60000/60000 [=====] - 44s 728us/step - loss: 0.0633 - acc: 0
Epoch 3/10
60000/60000 [=====] - 44s 728us/step - loss: 0.0406 - acc: 0
Epoch 4/10
60000/60000 [=====] - 44s 729us/step - loss: 0.0224 - acc: 0
Epoch 5/10
60000/60000 [=====] - 44s 729us/step - loss: 0.0141 - acc: 0
Epoch 6/10
60000/60000 [=====] - 44s 729us/step - loss: 0.0083 - acc: 0
Epoch 7/10
60000/60000 [=====] - 44s 729us/step - loss: 0.0044 - acc: 0
Epoch 8/10
60000/60000 [=====] - 44s 729us/step - loss: 0.0023 - acc: 0
Epoch 9/10
60000/60000 [=====] - 44s 729us/step - loss: 0.0015 - acc: 0
Epoch 10/10
60000/60000 [=====] - 44s 729us/step - loss: 0.0010 - acc: 1
Test loss: 0.021898437683159137
Test accuracy: 0.9925
```

▼ 7-layers

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(128, kernel_size=(2, 2), padding='same', activation='relu', input_shape=input_shape))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3), padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.5))
```



```
model.add(Conv2D(64, (3, 3),padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Conv2D(64, (5, 5),padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(32, (5, 5),padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Conv2D(32, (7, 7),padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(16, (7, 7),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(400, activation='relu'))
model.add(Dense(200, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam())

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_9"
```

Layer (type)	Output Shape	Param #
=====		
conv2d_33 (Conv2D)	(None, 28, 28, 128)	640
batch_normalization_11 (Batch Normalization)	(None, 28, 28, 128)	512
conv2d_34 (Conv2D)	(None, 28, 28, 128)	147584
max_pooling2d_18 (MaxPooling2D)	(None, 14, 14, 128)	0
dropout_13 (Dropout)	(None, 14, 14, 128)	0
conv2d_35 (Conv2D)	(None, 14, 14, 64)	73792
batch_normalization_12 (Batch Normalization)	(None, 14, 14, 64)	256
dropout_14 (Dropout)	(None, 14, 14, 64)	0
conv2d_36 (Conv2D)	(None, 14, 14, 64)	102464
batch_normalization_13 (Batch Normalization)	(None, 14, 14, 64)	256
dropout_15 (Dropout)	(None, 14, 14, 64)	0
max_pooling2d_19 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_37 (Conv2D)	(None, 7, 7, 32)	51232
batch_normalization_14 (Batch Normalization)	(None, 7, 7, 32)	128
dropout_16 (Dropout)	(None, 7, 7, 32)	0
conv2d_38 (Conv2D)	(None, 7, 7, 32)	50208
batch_normalization_15 (Batch Normalization)	(None, 7, 7, 32)	128
dropout_17 (Dropout)	(None, 7, 7, 32)	0
max_pooling2d_20 (MaxPooling2D)	(None, 3, 3, 32)	0
conv2d_39 (Conv2D)	(None, 3, 3, 16)	25104
max_pooling2d_21 (MaxPooling2D)	(None, 1, 1, 16)	0
flatten_9 (Flatten)	(None, 16)	0
dense_17 (Dense)	(None, 400)	6800
dense_18 (Dense)	(None, 200)	80200
dropout_18 (Dropout)	(None, 200)	0
dense_19 (Dense)	(None, 10)	2010
=====		
Total params: 541,314		

Trainable params: 540,674
 Non-trainable params: 640

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/12
60000/60000 [=====] - 44s 732us/step - loss: 0.5045 - acc: 0
Epoch 2/12
60000/60000 [=====] - 41s 680us/step - loss: 0.1028 - acc: 0
Epoch 3/12
60000/60000 [=====] - 41s 684us/step - loss: 0.0731 - acc: 0
Epoch 4/12
60000/60000 [=====] - 41s 683us/step - loss: 0.0622 - acc: 0
Epoch 5/12
60000/60000 [=====] - 41s 681us/step - loss: 0.0551 - acc: 0
Epoch 6/12
60000/60000 [=====] - 41s 681us/step - loss: 0.0519 - acc: 0
Epoch 7/12
60000/60000 [=====] - 41s 680us/step - loss: 0.0466 - acc: 0
Epoch 8/12
60000/60000 [=====] - 41s 682us/step - loss: 0.0440 - acc: 0
Epoch 9/12
60000/60000 [=====] - 41s 680us/step - loss: 0.0423 - acc: 0
Epoch 10/12
60000/60000 [=====] - 41s 681us/step - loss: 0.0387 - acc: 0
Epoch 11/12
60000/60000 [=====] - 41s 682us/step - loss: 0.0391 - acc: 0
Epoch 12/12
60000/60000 [=====] - 41s 680us/step - loss: 0.0372 - acc: 0
Test loss: 0.038856790525408
Test accuracy: 0.99
```

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
```

```

    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(256, kernel_size=(2, 2),padding='same',activation='relu',input_shape=input_shape))
model.add(Conv2D(128, (3, 3),padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3),padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (5, 5),padding='same', activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(32, (5, 5),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(26, (7, 7),padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(400, activation='relu'))
model.add(Dense(200, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam())

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```



```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_10"
```

Layer (type)	Output Shape	Param #
conv2d_40 (Conv2D)	(None, 28, 28, 256)	1280
conv2d_41 (Conv2D)	(None, 28, 28, 128)	295040
batch_normalization_16 (Batch Normalization)	(None, 28, 28, 128)	512
conv2d_42 (Conv2D)	(None, 28, 28, 128)	147584
max_pooling2d_22 (MaxPooling2D)	(None, 14, 14, 128)	0
conv2d_43 (Conv2D)	(None, 14, 14, 64)	73792
batch_normalization_17 (Batch Normalization)	(None, 14, 14, 64)	256
conv2d_44 (Conv2D)	(None, 14, 14, 32)	51232
batch_normalization_18 (Batch Normalization)	(None, 14, 14, 32)	128
conv2d_45 (Conv2D)	(None, 14, 14, 32)	25632
max_pooling2d_23 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_46 (Conv2D)	(None, 7, 7, 26)	40794
max_pooling2d_24 (MaxPooling2D)	(None, 3, 3, 26)	0
flatten_10 (Flatten)	(None, 234)	0
dense_20 (Dense)	(None, 400)	94000
dense_21 (Dense)	(None, 200)	80200
dropout_19 (Dropout)	(None, 200)	0
dense_22 (Dense)	(None, 10)	2010

```
=====  
Total params: 812,460  
Trainable params: 812,012  
Non-trainable params: 448
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/12  
60000/60000 [=====] - 63s 1ms/step - loss: 0.1832 - acc: 0.9  
Epoch 2/12  
60000/60000 [=====] - 60s 997us/step - loss: 0.0513 - acc: 0  
Epoch 3/12  
60000/60000 [=====] - 60s 996us/step - loss: 0.0415 - acc: 0  
Epoch 4/12  
60000/60000 [=====] - 60s 997us/step - loss: 0.0297 - acc: 0  
Epoch 5/12  
60000/60000 [=====] - 60s 998us/step - loss: 0.0284 - acc: 0  
Epoch 6/12  
60000/60000 [=====] - 60s 996us/step - loss: 0.0219 - acc: 0
```

```

Epoch 7/12
60000/60000 [=====] - 60s 997us/step - loss: 0.0241 - acc: 0
Epoch 8/12
60000/60000 [=====] - 60s 995us/step - loss: 0.0188 - acc: 0
Epoch 9/12
60000/60000 [=====] - 60s 995us/step - loss: 0.0188 - acc: 0
Epoch 10/12
60000/60000 [=====] - 60s 996us/step - loss: 0.0168 - acc: 0
Epoch 11/12
60000/60000 [=====] - 60s 997us/step - loss: 0.0152 - acc: 0
Epoch 12/12
60000/60000 [=====] - 60s 994us/step - loss: 0.0144 - acc: 0
Test loss: 0.03264233518390265
Test accuracy: 0.9923

```

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

```

```
# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(256, kernel_size=(2, 2),padding='same',activation='sigmoid',input_shape=i
model.add(Conv2D(128, (3, 3),padding='same', activation='sigmoid'))
model.add(Dropout(0.5))
model.add(Conv2D(128, (3, 3),padding='same', activation='sigmoid'))
model.add(Conv2D(128, (5, 5),padding='same', activation='sigmoid'))
model.add(Dropout(0.5))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (5, 5),padding='same', activation='sigmoid'))
model.add(Conv2D(128, (7, 7),padding='same', activation='sigmoid'))
model.add(Dropout(0.5))
model.add(Conv2D(128, (7, 7),padding='same', activation='sigmoid'))
model.add(Dropout(0.5))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(400, activation='sigmoid'))
model.add(Dense(200, activation='sigmoid'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adagra

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_11"
```

Layer (type)	Output Shape	Param #
conv2d_47 (Conv2D)	(None, 28, 28, 256)	1280
conv2d_48 (Conv2D)	(None, 28, 28, 128)	295040
dropout_20 (Dropout)	(None, 28, 28, 128)	0
conv2d_49 (Conv2D)	(None, 28, 28, 128)	147584
conv2d_50 (Conv2D)	(None, 28, 28, 128)	409728
dropout_21 (Dropout)	(None, 28, 28, 128)	0
max_pooling2d_25 (MaxPooling)	(None, 14, 14, 128)	0
conv2d_51 (Conv2D)	(None, 14, 14, 128)	409728
conv2d_52 (Conv2D)	(None, 14, 14, 128)	802944
dropout_22 (Dropout)	(None, 14, 14, 128)	0
conv2d_53 (Conv2D)	(None, 14, 14, 128)	802944
dropout_23 (Dropout)	(None, 14, 14, 128)	0
max_pooling2d_26 (MaxPooling)	(None, 7, 7, 128)	0
flatten_11 (Flatten)	(None, 6272)	0
dense_23 (Dense)	(None, 400)	2509200
dense_24 (Dense)	(None, 200)	80200
dropout_24 (Dropout)	(None, 200)	0
dense_25 (Dense)	(None, 10)	2010
Total params: 5,460,658		
Trainable params: 5,460,658		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 103s 2ms/step - loss: 2.3165 - acc: 0.

Epoch 2/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3020 - acc: 0.1

Epoch 3/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3017 - acc: 0.1

Epoch 4/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3015 - acc: 0.1

Epoch 5/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3017 - acc: 0.1

Epoch 6/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3014 - acc: 0.1


```

Epoch 7/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3015 - acc: 0.1
Epoch 8/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3015 - acc: 0.1
Epoch 9/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3013 - acc: 0.1
Epoch 10/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3014 - acc: 0.1
Epoch 11/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3014 - acc: 0.1
Epoch 12/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3014 - acc: 0.1
Test loss: 2.3010334442138674
Test accuracy: 0.1135

```

Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

```

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

print(x_train.shape)

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

```

```
# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(512, kernel_size=(2, 2),padding='same',activation='tanh',input_shape=inputs[0].shape))
model.add(Conv2D(256, (3, 3),padding='same', activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3),padding='same', activation='tanh'))
model.add(Conv2D(128, (5, 5),padding='same', activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (5, 5),padding='same', activation='tanh'))
model.add(Conv2D(64, (5, 5),padding='same', activation='tanh'))
model.add(Conv2D(32, (7, 7),padding='same', activation='tanh'))
model.add(Conv2D(16, (7, 7),padding='same', activation='tanh'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(400, activation='tanh'))
model.add(Dense(200, activation='tanh'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam)

model.fit(x_train, y_train,batch_size=batch_size,epochs=epochs,verbose=1,validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```



```
(60000, 28, 28)
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Model: "sequential_12"
```

Layer (type)	Output Shape	Param #
conv2d_54 (Conv2D)	(None, 28, 28, 512)	2560
conv2d_55 (Conv2D)	(None, 28, 28, 256)	1179904
max_pooling2d_27 (MaxPooling)	(None, 14, 14, 256)	0
conv2d_56 (Conv2D)	(None, 14, 14, 128)	295040
conv2d_57 (Conv2D)	(None, 14, 14, 128)	409728
max_pooling2d_28 (MaxPooling)	(None, 7, 7, 128)	0
conv2d_58 (Conv2D)	(None, 7, 7, 128)	409728
conv2d_59 (Conv2D)	(None, 7, 7, 64)	204864
conv2d_60 (Conv2D)	(None, 7, 7, 32)	100384
conv2d_61 (Conv2D)	(None, 7, 7, 16)	25104
max_pooling2d_29 (MaxPooling)	(None, 3, 3, 16)	0
flatten_12 (Flatten)	(None, 144)	0
dense_26 (Dense)	(None, 400)	58000
dense_27 (Dense)	(None, 200)	80200
dropout_25 (Dropout)	(None, 200)	0
dense_28 (Dense)	(None, 10)	2010
Total params: 2,767,522		
Trainable params: 2,767,522		
Non-trainable params: 0		

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 105s 2ms/step - loss: 2.4997 - acc: 0.

Epoch 2/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3530 - acc: 0.0

Epoch 3/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3278 - acc: 0.1

Epoch 4/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3145 - acc: 0.1

Epoch 5/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3095 - acc: 0.1

Epoch 6/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3083 - acc: 0.1

Epoch 7/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3080 - acc: 0.1

Epoch 8/12

60000/60000 [=====] - 99s 2ms/step - loss: 2.3074 - acc: 0.1

```

Epoch 9/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3076 - acc: 0.1
Epoch 10/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3069 - acc: 0.1
Epoch 11/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3072 - acc: 0.1
Epoch 12/12
60000/60000 [=====] - 99s 2ms/step - loss: 2.3066 - acc: 0.1
Test loss: 2.305227696228027
Test accuracy: 0.1135

```

▼ Summary

```

#pretty table
#c-Convolved layer M-Maxpool layer
from prettytable import PrettyTable
x = PrettyTable()

x.field_names = ["No of concoluted layers", "Batch_Normalization", "Activation_function", "

x.add_row([3, "Y", "Relu", "Adam", "Glorot Normal", "CCMCM", "Y", 0.9919])
x.add_row([3, "Y", "sigmoid", "Adagrad", "He_Normal", "CMCCM", "N", 0.9930])
x.add_row([3, "N", "tanh", "Adadelta", "Random Normal", "CCCM", "Y", 0.9904])
x.add_row([3, "N", "Relu", "sgd", "Glorot Normal", "CCC", "N", 0.9790])

x.add_row([5, "Y", "Relu", "Adam", "Random Normal", "CCMCCMCM", "Y", 0.9921])
x.add_row([5, "Y", "Relu", "Adam", "Random Normal", "CMCMCMCMCM", "N", 0.9901])
x.add_row([5, "N", "sigmoid", "Adagrad", "He Normal", "CCCMCCM", "Y", 0.1135])
x.add_row([5, "N", "tanh", "adadelta", "Glorot Normal", "CCCMCMCM", "N", 0.9925])

x.add_row([7, "Y", "Relu", "Adam", "Random Normal", "CCMCCMCCMCM", "Y", 0.9900])
x.add_row([7, "Y", "Relu", "Adam", "Random Normal", "CCCMCCCMCM", "N", 0.9923])
x.add_row([7, "N", "sigmoid", "Adagrad", "He Normal", "CCCMCCCM", "Y", 0.1135])
x.add_row([7, "N", "tanh", "adadelta", "Glorot Normal", "CMCCMCCCM", "N", 0.1135])

print(x)

```



No of concoluted layers	Batch_Normalization	Activation_function	optimizers
3	Y	Relu	Adam
3	Y	sigmoid	Adagrad
3	N	tanh	Adadelata
3	N	Relu	sgd
5	Y	Relu	Adam
5	Y	Relu	Adam
5	N	sigmoid	Adagrad
5	N	tanh	adadelata
7	Y	Relu	Adam
7	Y	Relu	Adam
7	N	sigmoid	Adagrad
7	N	tanh	adadelata

1.3 convoluted layers with sigmoid activation unit gives best accuarcy score followed by 7 convoluted layers wit
 2.5 convoluted layers with sigmoid activation function and adagrad optimizer and 7 conoluted layers with sigmo